

The Berkeley Sockets API

Networked Systems Architecture 3
Lecture 4




UNIVERSITY
of
GLASGOW

The Berkeley Sockets API


- Widely used low-level C networking API
- First introduced in 4.3BSD Unix
 - Now available on most platforms: Linux, MacOS X, Windows, FreeBSD, Solaris, etc.
 - Largely compatible cross-platform

Concepts



- Sockets provide a standard interface between network and application
- Two types of socket:
 - Stream – provides a virtual circuit service
 - Datagram – delivers individual packets
- Independent of network type:
 - Commonly used with TCP/IP and UDP/IP, but not specific to the Internet protocols
 - Only discuss TCP/IP sockets today

TCP/IP Connection



```
int fd = socket(...)
```

```
connect(fd, ..., ...)
```

```
write(fd, data, datalen)
```

```
read(fd, buffer, buflen)
```

```
close(fd)
```

```
int fd = socket(...)
```

```
bind(fd, ..., ...)
```

```
listen(fd, ...)
```

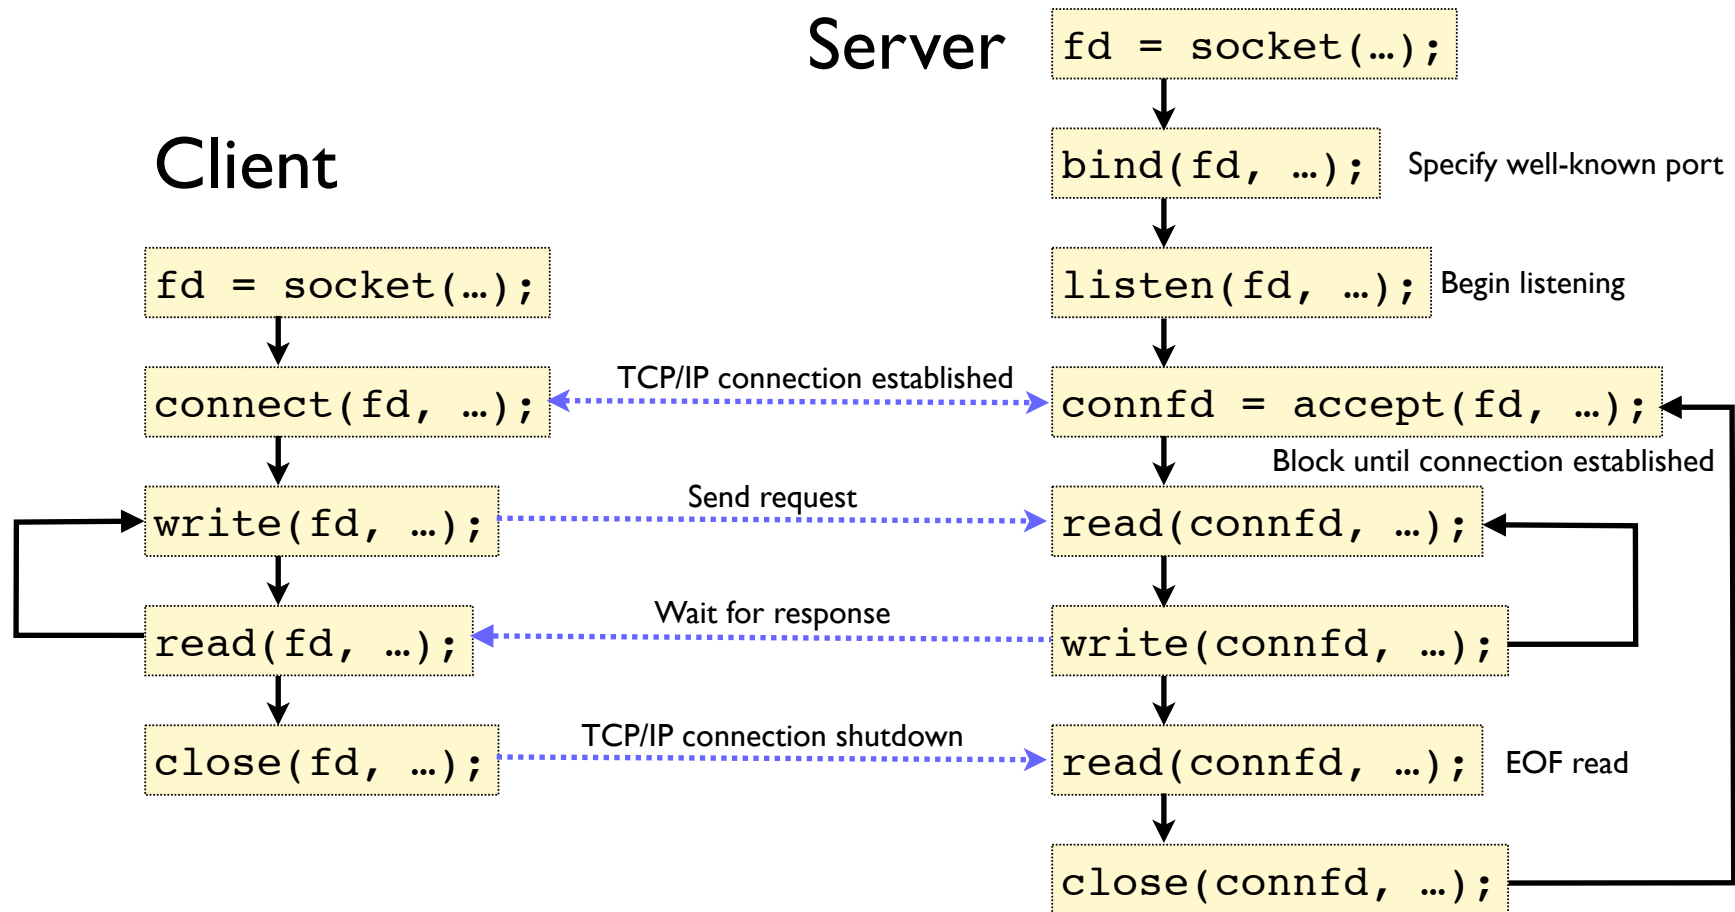
```
connfd = accept(fd, ...) ←
```

```
read(connfd, buffer, buflen)
```

```
write(connfd, data, datalen)
```

```
close(connfd)
```

TCP/IP Connection



Creating a socket

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int fd;
...
fd = socket(family, type, protocol);
if (fd == -1) {
    // Error: unable to create socket
    ...
}
...
```

AF_INET for IPv4

AF_INET6 for IPv6

SOCK_STREAM for TCP

SOCK_DGRAM for UDP

0 (not used for Internet sockets)

Create an unbound socket, not connected to network;
can be used as either a client or a server

Handling Errors

Socket functions return `-1` and set the global variable `errno` on failure

```
fd = socket(family, type, protocol);
if (fd == -1) {
    switch (errno) {
        case EPROTONOTSUPPORT :
            // Protocol not supported
            ...
        case EACCESS:
            // Permission denied
            ...
        default:
            // Other error...
            ...
    }
}
```

The Unix man pages should list the possible errors that can occur for each function

E.g. do “`man socket`” and see the **ERRORS** section

Binding a Server Socket

- Bind a socket to a port on a network interface
 - Needed to run servers on a well-known port
 - Not generally used on clients, since typically don't care which port used

```
#include <sys/types.h>
#include <sys/socket.h>

...
if (bind(fd, addr, addrlen) == -1) {
    // Error: unable to bind
    ...
}
...
```


Listening for Connections

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
if (listen(fd, backlog) == -1) {
    // Error
    . . .
}
. . .
```

Tell the socket to listen for new connections

The *backlog* is the maximum number of connections the socket will queue up, each waiting to be `accept ()`'ed

Connecting to a Server

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
if (connect(fd, addr, addrlen) == -1) {
    // Error: unable to open connection
    ...
}
...
```

Pointer to a struct sockaddr

Size of the struct in bytes

Tries to open a connection to the server
Times out after 75 seconds if no response

Specifying Addresses & Ports

- Must specify the address and port when calling `bind()` or `connect()`
 - The address can be either IPv4 or IPv6
 - Could be modelled in C as a union, but the designers of the sockets API chose to use a number of structs, and abuse casting instead

struct sockaddr

- Addresses specified using struct sockaddr

- Has a data field big enough to hold the largest address of any family
- Plus sa_len and sa_family to specify the length and type of the address
- Treats the address as an opaque binary string

```
struct sockaddr {  
    uint8_t      sa_len;  
    sa_family_t  sa_family;  
    char         sa_data[22];  
};
```

struct sockaddr_in

- Two variations exist for IPv4 and IPv6 addresses
 - Use struct sockaddr_in to hold an IPv4 address
 - Has the same size and memory layout as struct sockaddr, but interprets the bits differently to give structure to the address

```
struct in_addr {
    in_addr_t    s_addr;
};

struct sockaddr_in {
    uint8_t      sin_len;
    sa_family_t  sin_family;
    in_port_t    sin_port;
    struct in_addr sin_addr;
    char         sin_pad[16];
};
```

struct sockaddr_in6

- Two variations exist for IPv4 and IPv6 addresses
 - Use struct sockaddr_in6 to hold an IPv6 address
 - Has the same size and memory layout as struct sockaddr, but interprets the bits differently to give structure to the address

```
struct in6_addr {  
    uint8_t      s6_addr[16];  
};  
  
struct sockaddr_in6 {  
    uint8_t      sin6_len;  
    sa_family_t  sin6_family;  
    in_port_t    sin6_port;  
    uint32_t     sin6_flowinfo;  
    struct in6_addr sin6_addr;  
};
```

Working with Addresses

- Work with either `struct sockaddr_in` or `struct sockaddr_in6`
- Cast it to a `struct sockaddr` before calling the socket routines

```
struct sockaddr_in  addr;  
...  
// Fill in addr here  
...  
if (bind(fd, (struct sockaddr *) &addr, sizeof(addr)) == -1) {  
    ...  
}
```

Creating an Address: Manually

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
```

inet_pton() to convert address
htons() to convert port

```
struct sockaddr_in  addr;
...
inet_pton(AF_INET, "130.209.240.1", &addr.sin_addr);
addr.sin_family = AF_INET;
addr.sin_port   = htons(80);

if (connect(fd, (struct sockaddr *)&addr, sizeof(addr)) == -1) {
    ...
}
```


Creating an Address: DNS

- Prefer using DNS names to raw IP addresses
 - Use `getaddrinfo()` to look-up name in the DNS
 - Returns a linked list of `struct addrinfo` values, representing the addresses of the host

```
struct addrinfo {  
    int             ai_flags;        // input flags  
    int             ai_family;       // AF_INET, AF_INET6, ...  
    int             ai_socktype;     // IPPROTO_TCP, IPPROTO_UDP  
    int             ai_protocol;     // SOCK_STREAM, SOCK_DGRAM, ...  
    socklen_t       ai_addrlen;      // length of socket-address  
    struct sockaddr *ai_addr;         // socket-address for socket  
    char            *ai_canonname;   // canonical name of host  
    struct addrinfo *ai_next;        // pointer to next in list  
};
```

Connecting via a DNS Query

```
struct addrinfo          hints, *ai, *ai0;

memset(&hints, 0, sizeof(hints));
hints.ai_family      = PF_UNSPEC;
hints.ai_socktype    = SOCK_STREAM;
if ((i = getaddrinfo("www.google.com", "80", &hints, &ai0)) != 0) {
    printf("Unable to look up IP address: %s", gai_strerror(i));
    ...
}

for (ai = ai0; ai != NULL; ai = ai->ai_next) {
    fd = socket(ai->ai_family, ai->ai_socktype, ai->ai_protocol);
    if (fd == -1) {
        perror("Unable to create socket");
        continue;
    }

    if (connect(fd, ai->ai_addr, ai->ai_addrlen) == -1) {
        perror("Unable to connect");
        close(fd);
        continue;
    }
    ...
}
```

Accepting Connections

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int                connfd;
struct sockaddr_in cliaddr;
socklen_t          cliaddrlen = sizeof(cliaddr);
...
connfd = accept(fd, &cliaddr, &cliaddrlen);
if (connfd == -1) {
    // Error
    ...
}
...
```

Accept a connection, returning a new file descriptor for that connection (`connfd`) and the client's address (`cliaddr`)

Accepting Connections

- A TCP/IP server may have multiple connections outstanding
 - Can `accept ()` connections one at a time, handling each request in series
 - Can `accept ()` connections and start a new thread for each, allowing it to process several in parallel
- Each call to `accept ()` returns a new file descriptor for the new connection

Reading and Writing Data

```
#define BUFLen 1500
...
ssize_t i;
ssize_t rcount;
char    buf[BUFLen];
...
rcount = read(fd, buf, BUFLen);
if (rcount == -1) {
    // Error has occurred
    ...
}
...
for (i = 0; i < rcount; i++) {
    printf("%c", buf[i]);
}
```

Read up to BUFLen bytes of data from connection. Blocks until data available to read.

Returns actual number of bytes read, or `-1` on error.

Data read from the connection is *not* null terminated.

Reading and Writing Data

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main()
{
    char    x[] = "Hello, world!";
    char    *y   = malloc(14);

    sprintf(y, "Hello, world!");

    printf("x = %s\n", x);
    printf("y = %s\n", y);

    printf("sizeof(x) = %d\n", sizeof(x));
    printf("sizeof(y) = %d\n", sizeof(y));

    printf("strlen(x) = %d\n", strlen(x));
    printf("strlen(y) = %d\n", strlen(y));

    return 0;
}
```

What gets printed?

Why?

Reading and Writing Data

```
char data = "Hello, world!";  
int  datalen = strlen(data);  
...  
if (write(fd, data, datalen) == -1) {  
    // Error has occurred  
    ...  
}  
...
```

Send data on a TCP/IP connection. Blocks until all data can be written.

Returns actual number of bytes written, or -1 on error.

Closing a Socket

```
#include <unistd.h>
```

```
close(fd);
```

Close and destroy a socket

Close the file descriptor for each connection, then the file descriptor for the underlying socket

Questions?